

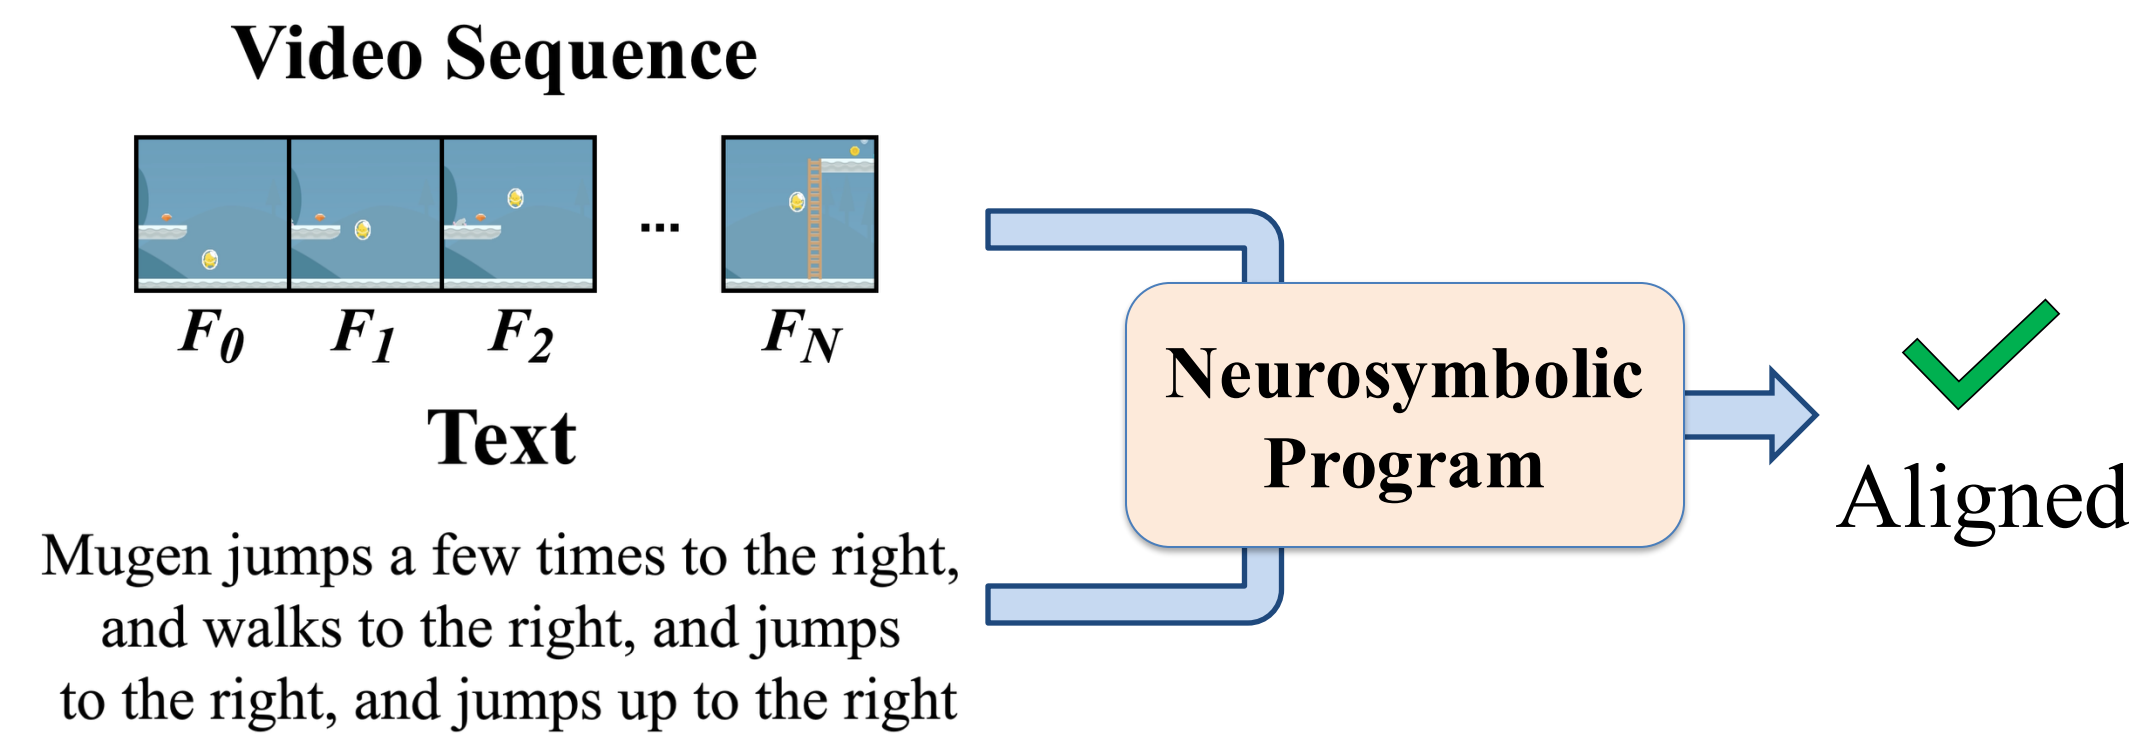
Dolphin: A Programmable Framework for Scalable Neurosymbolic Learning

Aaditya Naik, Jason Liu, Claire Wang, Amish Sethi, Saikat Dutta, Mayur Naik, Eric Wong

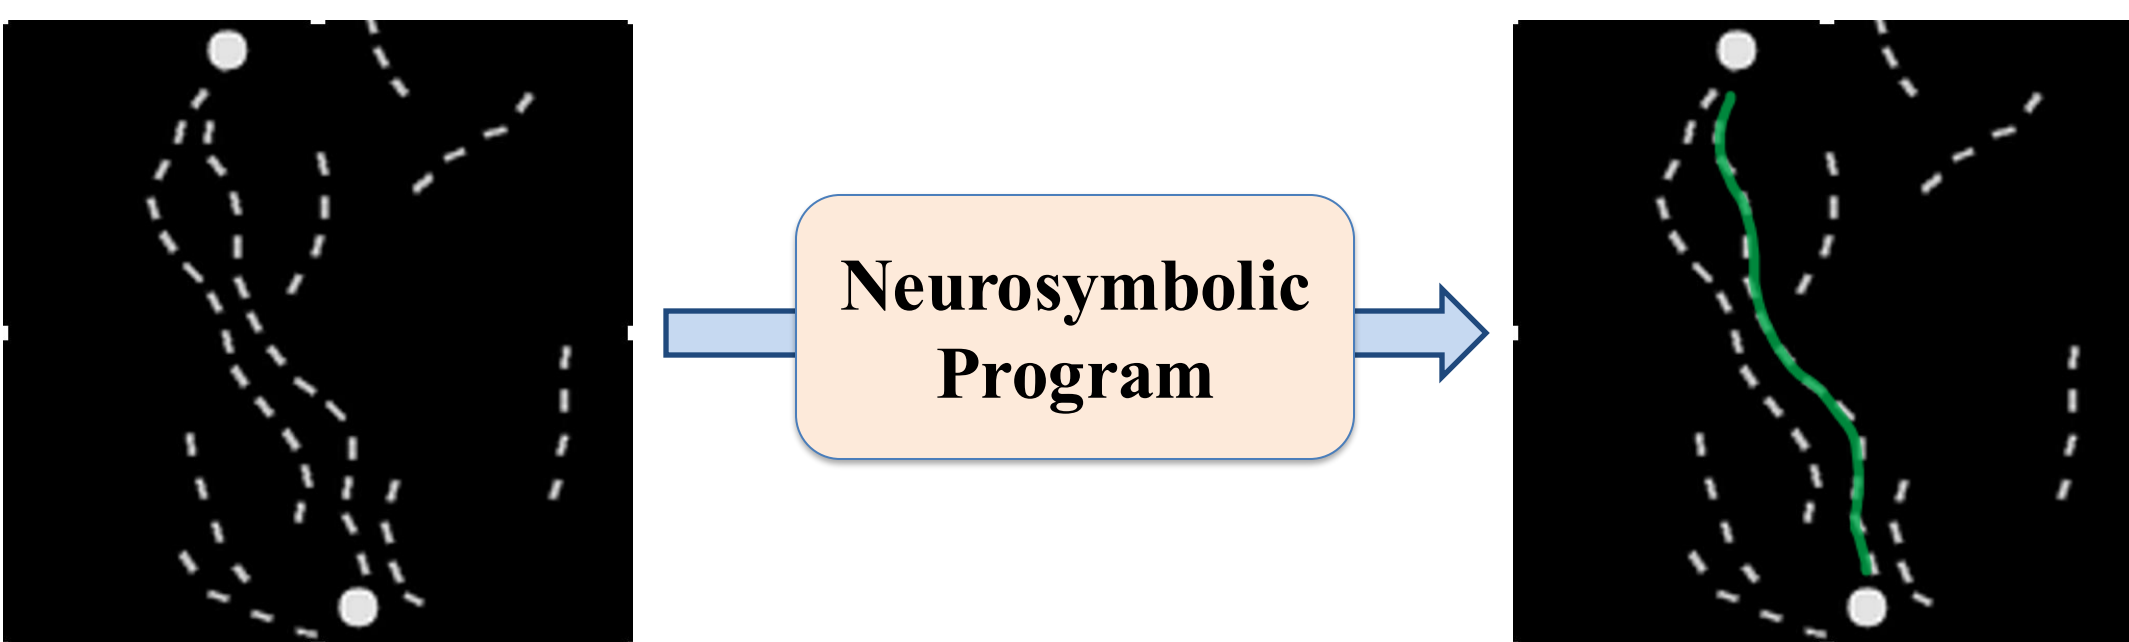
Motivation: Scaling Neurosymbolic Programming

Task

Video-Text Alignment


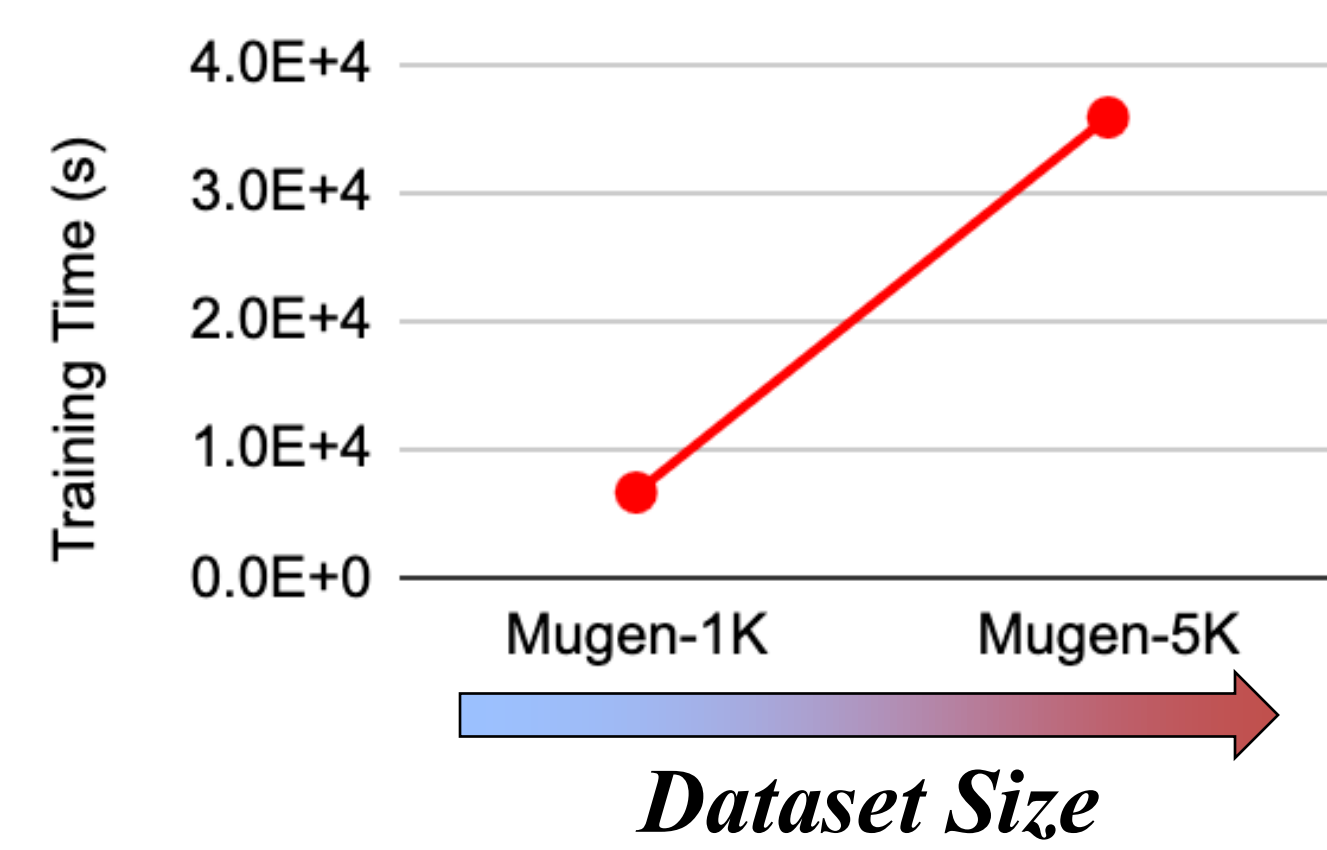


Visual Path-Finding



Training Time

with Scallop



An Example Dolphin Program

Task: Add N MNIST Digits

```
class SumNNNet(torch.nn.Module):
    def __init__(self):
        super(SumNNNet, self).__init__()
        self.CNN = MNISTNet()

    def forward(self, imgs):
        d = range(10)
        D_res = Distribution(self.CNN(imgs[0]), d)
        for i in range(1, len(imgs)):
            D_i = Distribution(self.CNN(imgs[i]), d)
            D_res = apply(D_res, D_i, lambda x,y: x + y)
        return get_logits(D_res)
```

Tight integration
with PyTorch

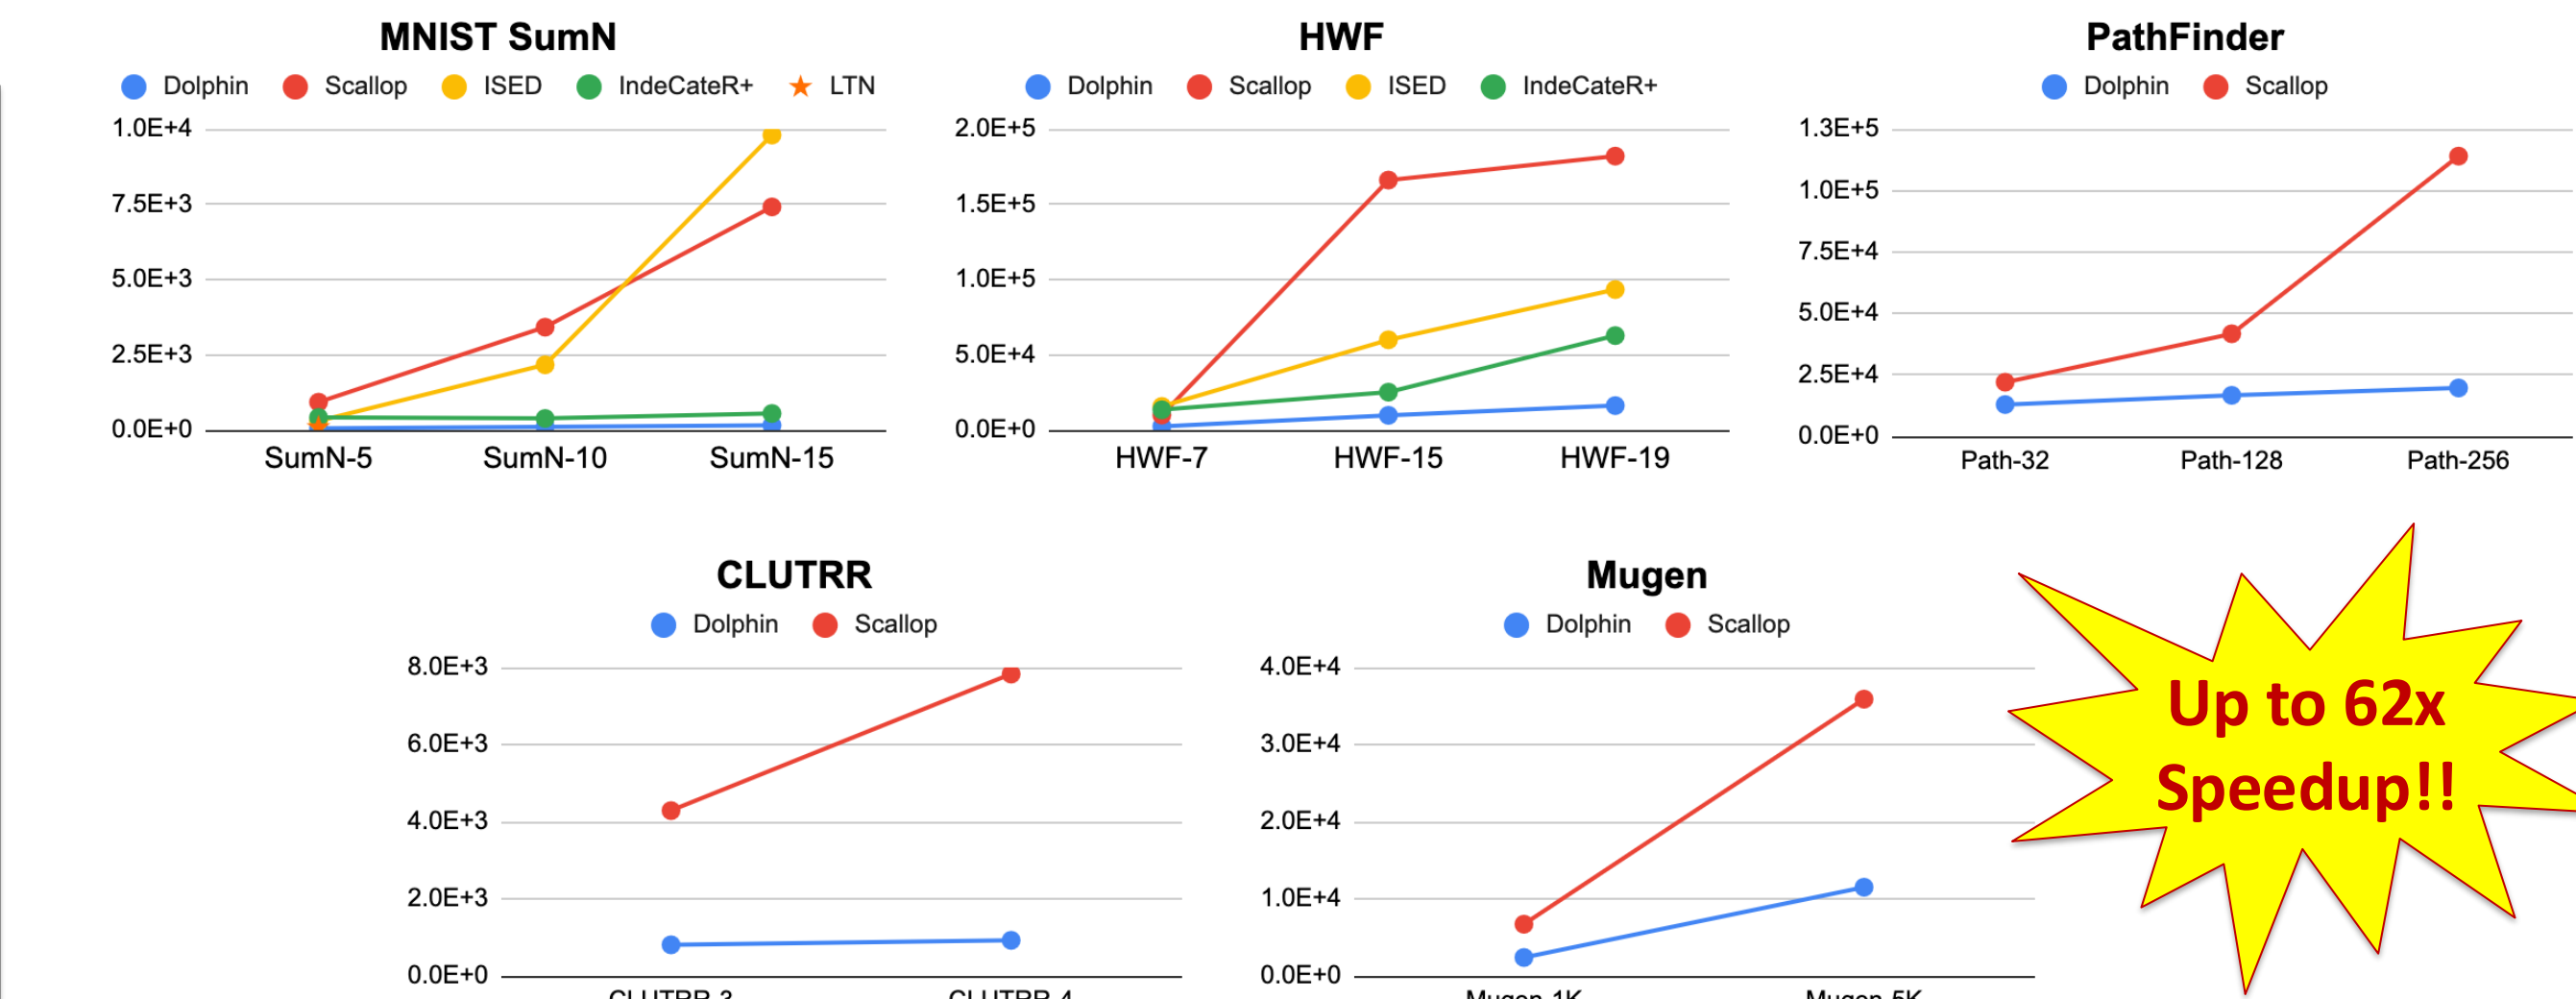
Intuitive Pythonic
Primitives

Recursion, Control-Flow,
Blackbox Functions

GPU-driven Vectorized
Probabilistic Computations


Experimental Results

Dolphin Speeds up Training



Dolphin Training Scales Better as Complexity Grows

Dolphin achieves SOTA Accuracies



Dolphin Splits Execution over CPU and GPU to allow Scaling

